

IN THIS CHAPTER

- » Exploring the basic principles of service-oriented apps
- » Creating service-oriented apps
- » Making XML Web services available

Online Chapter **4**

Getting Acquainted with Web Services

A *web service* is the provision of functionality over the Internet using an open interface called an Application Programming Interface (API). A web page provides functionality that you can see; a web service provides access to the underlying data or other resources using method calls, in a format that you can use in another application. Web services are a kind of web-based resource.

Web services are straightforward for the most part, but can provide a complex set of services underneath. For example, some web services now provide a complete computing environment that can replace your existing in-house environment. Web services are generally standards driven, just as HTML is, and the World Wide Web Consortium (W3C) owns its documentation. Web services have been a hot topic for quite some time, but only in the past ten years or so (in step with the ubiquitous nature of the Internet) have they become a viable option for the delivery of hard-to-find software functionality. A few different web services formats exist in the .NET world. The web services discussed in this chapter solve two basic problems:

- » Making part of your application code and underlying data available past the physical boundary of the application.

» Making a distributed middle to your application so that you can *scale up* (add computing, memory, data, or other resources) if your site suddenly has a lot of traffic.



REMEMBER

Web services perform all sorts of tasks today. Some of them are immense, such as Amazon Web Services (AWS). (See *AWS For Admins For Dummies* and *AWS For Developers For Dummies*, by John Paul Mueller [Wiley] for details.) This chapter provides you with some code after covering a few basic web service principles, but you should consider this coverage just the tip of the iceberg.

Understanding Web Services

Web services give you a way to extend methods past the normal boundary of a software system. You usually write something like the following chunk of code to start building a program:

```
public bool AddStuff(String stuff)
{
    //Add it here
    return true;
}
```

and then you call it like this:

```
bool DidItWork = AddStuff("This is the new Stuff");
```

Using Simple Object Access Protocol (SOAP) web services, however, you call the function this way:

```
POST /Service1.asmx HTTP/1.1
Host: localhost
Content-Type: application/soap+xml; charset=utf-8
Content-Length: length
<?xml version="1.0" encoding="utf-8"?>
<soap12:Envelope xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
    xmlns:xsd="http://www.w3.org/2001/XMLSchema" xmlns:soap12="http://www.w3.
    org/2003/05/soap-envelope">
  <soap12:Body>
    <AddStuff xmlns="http://tempuri.org/">
      <stuff>string</stuff>
    </AddStuff>
  </soap12:Body>
</soap12:Envelope>
```

and then see a response like this one:

```
HTTP/1.1 200 OK
Content-Type: application/soap+xml; charset=utf-8
Content-Length: length
<?xml version="1.0" encoding="utf-8"?>
<soap12:Envelope xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
  xmlns:xsd="http://www.w3.org/2001/XMLSchema" xmlns:soap12="http://www.w3.
  org/2003/05/soap-envelope">
  <soap12:Body>
    <AddStuffResponse xmlns="http://tempuri.org/">
      <AddStuffResult>boolean</AddStuffResult>
    </AddStuffResponse>
  </soap12:Body>
</soap12:Envelope>
```

You might wonder why anyone would want to build a function that way. The section entitled “Building Service-Oriented Applications,” later in this chapter, answers your questions. For now, here are some concepts that make building web services easier to handle. The following sections discuss the three basic principles that define SOAP web services:

- » **Loosely coupled:** They don’t require a constant connection to the server.
- » **Contract driven:** They provide an interface that describes all their functionality.
- » **More likely to be chunky, not chatty:** Rather than lots of properties with single values, they provide big methods that return collections.

Loosely coupled

Because web services are, like web applications, *loosely coupled*, web service conversations aren’t guaranteed to make sense. You might get a sentence from three minutes ago, after four other sentences have gone past, or you might not hear from them again after the first line.

When you stop to think about it, loose coupling makes sense: Web service calls are just like navigation in a web page. Sometimes you click a link and then close your browser, and sometimes you click the link twice. Web applications are up to the whim of the user, and so are web services.

For this reason, a client of a web service must be *loosely coupled* to the service. For example, you don’t want to make a user wait until a web service call is complete. You must call a web service method asynchronously and have the result show up when it is ready.

KINDS OF WEB SERVICES

If you have spent much time reviewing web services, you know that Microsoft initially started with a strong focus on SOAP because it mimics the functionality provided by the Component Object Model (COM) used for many elements in Windows. In addition, SOAP is highly structured and extensible. Unfortunately, it's not very flexible, consumes a large number of resources, and is prone to problematic errors. In many respects, Microsoft products still have a strong SOAP focus despite these issues, which is why you see SOAP covered in this chapter.

The rest of the world has moved on to REpresentational State Transfer (REST), which is software architecture, instead of a specific implementation covered by standards. REST makes it significantly easier to exchange data between disparate systems and allows more than one way to perform that data exchange. Large web services, such as AWS and Google's various web services, rely heavily on REST. You see REST covered in Online Chapter 7. If you want brief coverage of REST now, check out the article at http://www.service-architecture.com/articles/web-services/representational_state_transfer_rest.html.

This book doesn't cover another major web service contender, JavaScript Object Notation (JSON). You see JSON used side by side with many REST implementations. For example, AWS uses JSON to define security policies and other configuration information. JSON provides an extremely simple method of transferring data using key/value pairs in dictionary-style structures. You can read more about JSON at http://www.service-architecture.com/articles/web-services/javascript_object_notation_json.html and find a tutorial for it at https://www.w3schools.com/js/js_json_intro.asp. It's possible to work with JSON in C# (see the article at <https://docs.microsoft.com/en-us/dotnet/framework/wcf/feature-details/how-to-serialize-and-deserialize-json-data>), but most developers today rely on other techniques to work with JSON that are outside the scope of this book.

Fortunately, .NET handles the asynchronous call. You just tell your application that you're calling a service asynchronously and then provide a delegate for the service to call when it's ready. As long as you handle the code properly, the service will work as expected. This example shows a potential implementation of the loose coupling with an asynchronous call (it's not executable code):

```
using System;
using System.Collections.Generic;
using System.Linq;
using System.Text;

namespace ConsoleApplication1
{
```

```

class Program
{
    static void Main(string[] args)
    {
        //Here YOU are calling the service asynchronously.
        AddStuffReference.Service1SoapClient client =
            new AddStuffReference.Service1SoapClient();
        client.AddStuffCompleted +=
            new EventHandler(client_AddStuffCompleted);
        client.AddStuffAsynch("This is the stuff");
    }

    //This method is called when the response comes back.
    //No timers or anything. .NET handles it for you.
    void proxy_AddStuffCompleted(object sender,
        AddStuffReference.Service1SoapClient.AddStuffCompletedEventArgs e)
    {
        string result = e.Result.ToString();
    }
}

```



REMEMBER

The preceding example only protects your user interface from experiencing a tie-up. You still have no real indication that the messages will ever be delivered, so you should never write software that depends on the delivery of the data from the service. It has to fail gracefully.

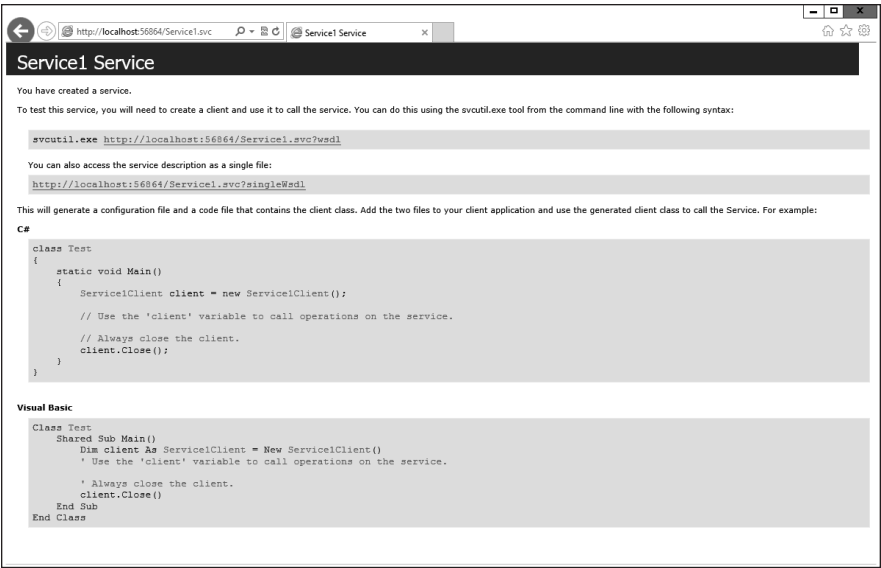
Contract driven

In client-server development, an interface defines a contract of sorts between the domain model and the user interface. You can use that contract to drive development in the same way that you use a contract to drive a business deal. If you know what a client requires and the client knows what you require, the two parties can quickly and easily make a deal.

A SOAP web service is required to have an interface that conforms to the Web Services Description Language (WSDL) standard. WSDL describes expected inputs and allowed outputs just as an interface would, creating a contract between the service provider and the client.

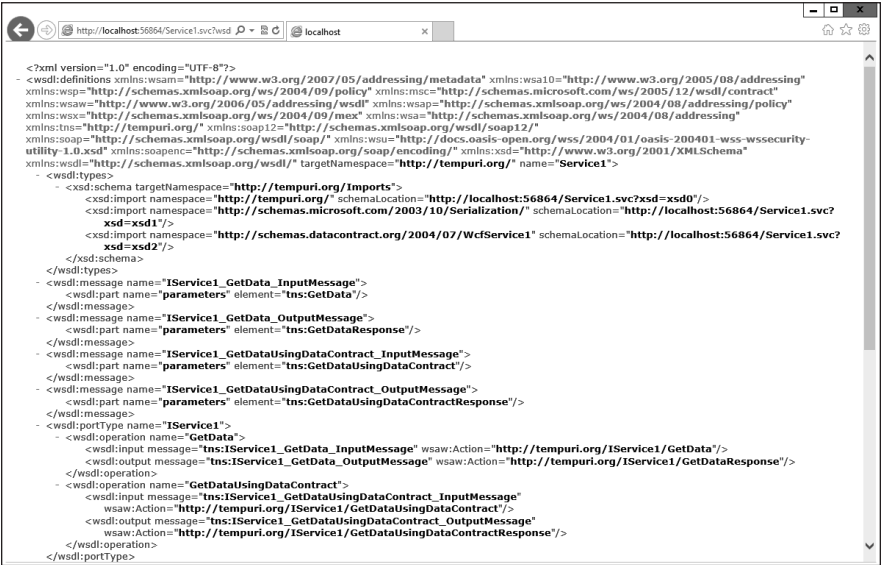
.NET creates the contract automatically from the code for the service. If you call a service in the browser (which isn't the way it's designed to be called) from the development machine, you see a test screen similar to the one shown in Figure OC4-1. Even though the browser gains some basic important information about the service, it isn't the contract. The contract is shown behind the Service Description link.

FIGURE OC4-1:
The default
service overview.



Clicking the Service Description link appends the text `?WSDL` to the URI, and the web browser shows the contract. The client system used this contract to determine exactly which information the service wants and how it will respond to input. Though the WSDL is displayed in a browser, as shown in Figure OC4-2, the interesting part is the XML that's behind the screen.

FIGURE OC4-2:
The Service
Description
in WSDL.



For compilable languages, the client machine does essentially the same thing no matter which platform it's on (Java or Basic or PHP or Ruby): At compile time, it reads the WSDL and creates a proxy that the client talks to. This

- » Brokers the communication process between the client and the service
- » Provides type safety, if it's supported
- » Generally makes your life easier because WSDL (and the contract it provides) is an important part of a service developer's work



REMEMBER

The service provider can change the service anytime it wants, and can even forgo updating the WSDL, which means that your application environment might not know the real requirements for making a call. However, .NET generates the WSDL automatically.

Chunky versus chatty

Although the chunky-versus-chatty services issue might sound like a face-off between candy bars, it isn't. Rather than perform small, incremental operations via web services, you use them to make large, sweeping "strokes." The length of individual calls might be larger, but the number of communications is fewer, which reduces network overhead.

Suppose that an application changes the settings of a piece of hardware located miles away from it — perhaps your home heating system (which is a good use of a service). The client application (your computer) is local, the remote device (your heater) is remote, and a network (probably the Internet) is in the middle. If you have a service with a series of individual controls, such as `TurnFanOn` and `TurnFanOff`, it's a common interface for a local application, such as connecting a heater directly to your computer. The following chunk of code gives an example of that chatty sort of interface:

```
namespace HomeHeater
{
    public class Chatty : IChatty
    {
        public bool TurnFanOn()
        {
            return true;
        }
        public bool TurnFanOff()
        {
            return true;
        }
    }
}
```

```

        public bool SetTemperature(int newTemperature)
        {
            return true;
        }
        public bool SetFanSpeed(string newFanSpeed)
        {
            return true;
        }
    }

```

The interface in this example is *chatty* (in case you couldn't tell from its class name). Every time you change a knob on the controller, a call is made to a service. Move the temperature to 72 degrees and call the service. Turn the fan to its High setting and call the service. Turn the temperature back down to 71 degrees and call the service. The client “chats” with the service.

Nothing is intrinsically wrong with this implementation. For a service, though, with more network overhead for every call, it isn't the best way to build methods. Instead, you want your client to change settings and then pull a big lever on the side to make all the changes at one time.

A *chunky* interface provides a domain model for the client to use, with properties to set. Then, after all settings are in the model object, you send the whole shootin' match to the service. It looks like this:

```

namespace HomeHeater
{
    public class Chunky : IChunky
    {
        public bool UpdateHeaterSettings(HeaterModule heaterModule)
        {
            return true;
        }
    }
    public class HeaterModule
    {
        public int Temperature { get; set; }
        public bool FanOn { get; set; }
        public string FanSpeed { get; set; }
    }
}

```

Now, no matter how often you make changes, the service is called only when you “pull the big lever” (update it). Whether this action prevents users from updating after every change depends on your interface. In general, this design principle is the best one for service development.

Building Service-Oriented Applications

The first and most obvious use of a service is to create a Service-Oriented Application, or SOA. This overloaded, overused term means that your application (like the heating system mentioned in the previous section) uses services to assist a remote client with server communications. The concept is simple: The user interface calls a service at some point in the process to communicate with the server. You do this for two reasons:

» **Scalability:** It's the main reason to use a service inside an application — especially a web application.

Up to a point, most web applications have built-in scaling. If you need more access points, you just add more servers and then use a device to sort the traffic to another machine, as shown in Figure OC4-3.




FIGURE OC4-3:
Scaling on a
simple website.

Web applications are different, though, because layers have different scalability needs. Sometimes the database is loaded, and sometimes the web server is. You must be able to separate the layers of the application, as shown in Figure OC4-4 — that's where services start to become useful.

Because the services use a common format — usually XML over HTTP — you can install parts of the application on their own machines to isolate them physically. Because the functionality of the deleted part is called via a service, you can scale the application horizontally — from one to multiple servers.




FIGURE OC4-4:
Separating
the layers.

- » **Reusability:** Every organization has a list of its participants — clients, users, voters, cooks, or mailing list subscribers, for example. Regardless of the type of participant, all the people on the list have first and last names and other identifying characteristics.
- » It seems that nearly every application written today has a table of People. Savvy programmers use slick tricks to keep these tables in sync or to share information, for example, but only one way exists to share the People table — by using a data silo, as shown in Figure OC4-5.




FIGURE OC4-5:
A data silo.

A data silo works this way:

1. A database on some server somewhere contains all participants' demographic information.
2. The database is surrounded by a service layer containing all allowed operations.
3. The service layer is consumed by all other applications in the system, as shown in Figure OC4-5.

This set of steps shows you the concept of reuse. It isn't about bits of code (no matter what you might hear); it's about *data*. Services help to provide access to the data silo.

Providing XML Web Services

A common use of services is to give other programmers public access to your information by using an open API that can be consumed anywhere. You might recall the WSDL file, described in the earlier section “Contract driven” — it can make your cool function or valuable data available to anyone who needs it.

In the demographic silo example (described in the “Building Service-Oriented Applications,” earlier in this chapter), if you have a valuable mailing list and you want to give your customers access to it, you can send them the list. If you do, however, they can use it forever. Instead, suppose that you could bill your customers every time they use the list. If you provide web services to implemented functions on the list, you can then track actual usage of the list.

You can see this concept, known as provision past the boundary, everywhere. Web services can be controlled as any other web application can. The field-level or operation-level security that used to be handled at the database level can now be built into a semipublic API.

Building Three Sample Apps

You can build services in C# in a few ways, as described in the Online Chapters 5–7. These chapters show you how to build *the same app* in three different ways. The app, the SHARP conference management system (or a simple version of it), was built using these formats:

» **ASMX:** This first version is in the venerable ASMX format: ASP.NET web services. Though this format has been largely superseded by WCF, it's still

available. Online Chapter 5 covers it because it's still a viable solution for certain situations.

- » **WCF:** Windows Communication Foundation (WCF), the most important service platform for all things Microsoft, is given enough screen time in Online Chapter 6 to get you started. The topic is broad, and you can find lots of references to other resources.
- » **REST:** REST, the method you likely find used for most web services today, appears in Online Chapter 7. Remember that REST isn't a protocol — it's an architectural style that specifies a way to build a web service rather than provide rules on its implementation. You could create a REST implementation that relies on SOAP just as easily as you can build one that relies on HTTP (which is the standard method). The issue of what REST represents precisely is so fraught with political positioning that this book doesn't really delve into the details. You can read an informed article on the topic at <https://www.upwork.com/hiring/development/soap-vs-rest-comparing-two-apis/>, however.

The entire process uses a single-entity data model, shown in Figure OC4-6, and provides allowed operations to it.

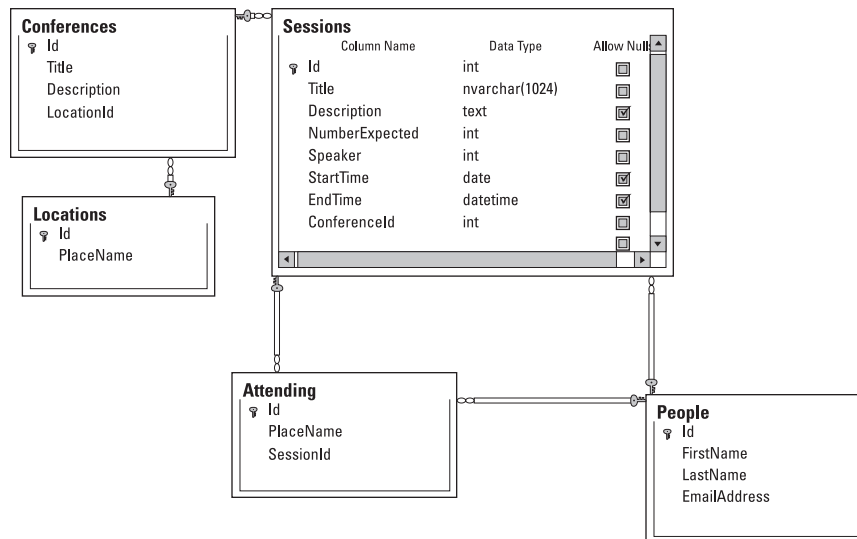


FIGURE OC4-6:
The simplified
SHARP data
model.

The idea is to provide a service model similar to the heating system platform described earlier in this chapter, except using a real, stateful, data-driven application. This chapter and the following three chapters show you how to solve three different problems. Then you can shoot for three different outcomes. Otherwise, what fun would it be?